

Regression

Logistic Regression

```
In [16]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.metrics import mean_squared_error

# Load the dataset
data = pd.read_csv('diabetes_csv.csv')

# Display the first few rows of the data to understand it
print(data.head())

# Select 'Age' as the feature (X) and 'Outcome' as the target (y)
X = data[['Age']] # Using double brackets to keep X as a DataFrame
y = data['Outcome']

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_st

# Train a logistic regression model using only 'Age' as the predictor
model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)

# Predict the outcomes for the test set
y_pred = model.predict(X_test)

# Predict the probabilities for the test set
y_pred_proba = model.predict_proba(X_test)[: , 1] # Predicted probabilities for the

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")

# Calculate precision
precision = precision_score(y_test, y_pred)
print(f"Precision: {precision:.4f}")

# Calculate recall
recall = recall_score(y_test, y_pred)
print(f"Recall: {recall:.4f}")

# Calculate F1-score
f1 = f1_score(y_test, y_pred)
print(f"F1 Score: {f1:.4f}")
```

```

# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=model.classes_)

# Plot confusion matrix
plt.figure(figsize=(8, 6))
disp.plot(cmap=plt.cm.Blues, values_format='d')
plt.title('Confusion Matrix')
plt.show()

# Calculate Mean Squared Error (MSE)
mse = mean_squared_error(y_test, y_pred_proba)
print(f"Mean Squared Error (MSE): {mse:.4f}")

# Calculate Root Mean Squared Error (RMSE)
rmse = np.sqrt(mse)
print(f"Root Mean Squared Error (RMSE): {rmse:.4f}")

# Sort values for a smooth line plot
X_test_sorted, y_pred_proba_sorted = zip(*sorted(zip(X_test['Age'], y_pred_proba)))

# Plotting Age vs Outcome using Logistic regression probabilities
plt.figure(figsize=(10, 6))
plt.plot(X_test_sorted, y_pred_proba_sorted, color='red', label='Logistic Regression')
plt.scatter(X_test['Age'], y_test, color='blue', alpha=0.5, label='Actual Data Points')
plt.xlabel('Age')
plt.ylabel('Outcome Probability')
plt.title('Logistic Regression: Age vs Outcome')
plt.legend()
plt.show()

```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1


Accuracy: 0.5885

Precision: 0.3529

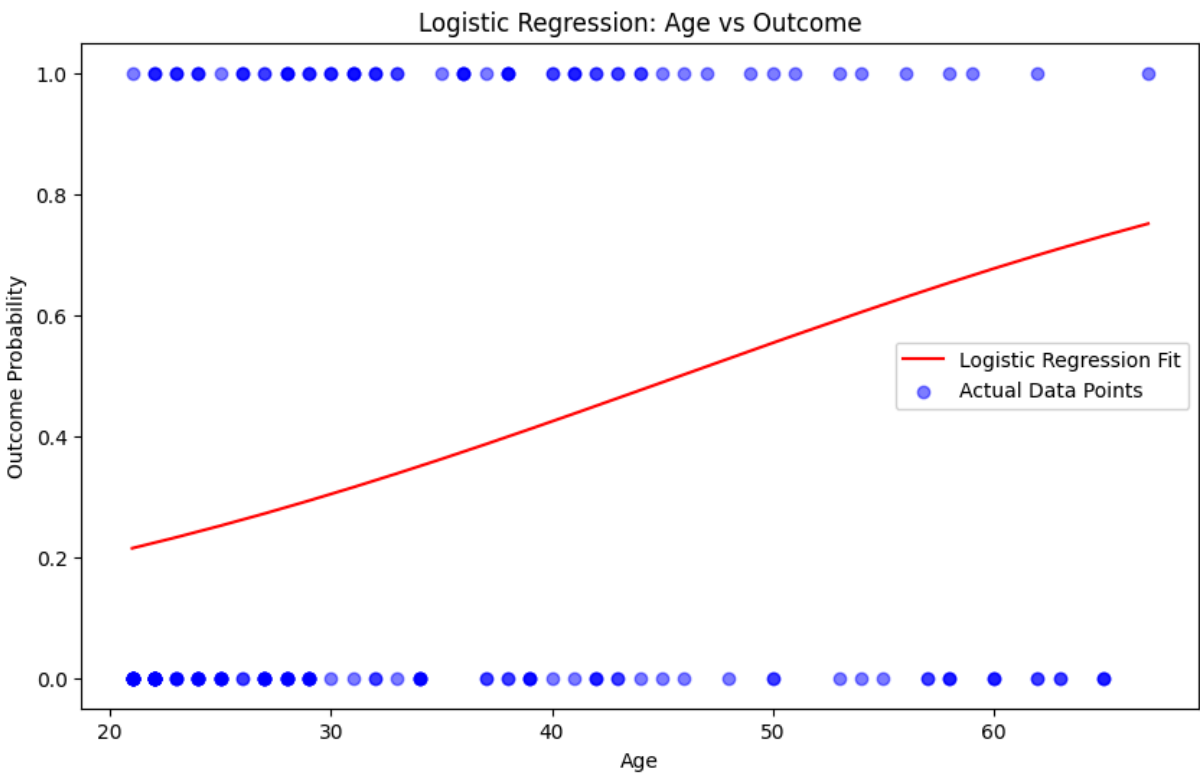
Recall: 0.1739

F1 Score: 0.2330

<Figure size 800x600 with 0 Axes>



Mean Squared Error (MSE): 0.2383
Root Mean Squared Error (RMSE): 0.4881



Linear Regression

```

In [17]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Load the dataset
data = pd.read_csv('diabetes_csv.csv')

# Display the first few rows of the data to understand it
print(data.head())

# Select 'Age' as the feature (X) and 'Outcome' as the target (y)
X = data[['Age']] # Using double brackets to keep X as a DataFrame
y = data['Outcome']

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_st

# Train a Linear regression model using only 'Age' as the predictor
model = LinearRegression()
model.fit(X_train, y_train)

# Predict the outcomes for the test set
y_pred = model.predict(X_test)

# Calculate Mean Squared Error (MSE)
mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error (MSE): {mse:.4f}")

# Calculate Root Mean Squared Error (RMSE)
rmse = np.sqrt(mse)
print(f"Root Mean Squared Error (RMSE): {rmse:.4f}")

# Calculate R-squared (R2) score
r2 = r2_score(y_test, y_pred)
print(f"R-squared (R2): {r2:.4f}")

# Plotting Age vs Outcome using Linear regression line
plt.figure(figsize=(10, 6))
plt.scatter(X_test['Age'], y_test, color='blue', alpha=0.5, label='Actual Data Point')
plt.plot(X_test['Age'], y_pred, color='red', linewidth=2, label='Linear Regression')
plt.xlabel('Age')
plt.ylabel('Outcome')
plt.title('Linear Regression: Age vs Outcome')
plt.legend()
plt.show()

```


	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

Mean Squared Error (MSE): 0.2362

Root Mean Squared Error (RMSE): 0.4860

R-squared (R2): -0.0258



SGD

SGD w/ Linear Regression

```
In [19]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import SGDRegressor
from sklearn.metrics import mean_squared_error, r2_score

# Load the dataset
data = pd.read_csv('diabetes_csv.csv')
```

```

# Display the first few rows of the data to understand it
print(data.head())

# Select 'Age' as the feature (X) and 'Outcome' as the target (y)
X = data[['Age']] # Using double brackets to keep X as a DataFrame
y = data['Outcome']

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_st

# Train a linear regression model using SGD
model = SGDRegressor(max_iter=1000, tol=1e-3, random_state=42)
model.fit(X_train, y_train)

# Predict the outcomes for the test set
y_pred = model.predict(X_test)

# Calculate Mean Squared Error (MSE)
mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error (MSE): {mse:.4f}")

# Calculate Root Mean Squared Error (RMSE)
rmse = np.sqrt(mse)
print(f"Root Mean Squared Error (RMSE): {rmse:.4f}")

# Calculate R-squared (R2) score
r2 = r2_score(y_test, y_pred)
print(f"R-squared (R2): {r2:.4f}")

# Plotting Age vs Outcome using linear regression line
plt.figure(figsize=(10, 6))
plt.scatter(X_test['Age'], y_test, color='blue', alpha=0.5, label='Actual Data Point')
plt.plot(X_test['Age'], y_pred, color='red', linewidth=2, label='Linear Regression')
plt.xlabel('Age')
plt.ylabel('Outcome')
plt.title('Linear Regression with SGD: Age vs Outcome')
plt.legend()
plt.show()

```


	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

Mean Squared Error (MSE): 0.3848

Root Mean Squared Error (RMSE): 0.6204

R-squared (R2): -0.6716



SGD w/ Logistic Regression

```
In [21]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import SGDClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.metrics import mean_squared_error

# Load the dataset
data = pd.read_csv('diabetes_csv.csv')

# Display the first few rows of the data to understand it
print(data.head())

# Select 'Age' as the feature (X) and 'Outcome' as the target (y)
X = data[['Age']] # Using double brackets to keep X as a DataFrame
y = data['Outcome']

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_st

# Train a logistic regression model using SGD
model = SGDClassifier(max_iter=1000, tol=1e-3, random_state=42, loss='log_loss')
model.fit(X_train, y_train)

# Predict the outcomes for the test set
y_pred = model.predict(X_test)

# Predict the probabilities for the test set using the decision function
```

```
y_pred_proba = model.decision_function(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")

# Calculate precision
precision = precision_score(y_test, y_pred)
print(f"Precision: {precision:.4f}")

# Calculate recall
recall = recall_score(y_test, y_pred)
print(f"Recall: {recall:.4f}")

# Calculate F1-score
f1 = f1_score(y_test, y_pred)
print(f"F1 Score: {f1:.4f}")

# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=model.classes_)

# Plot confusion matrix
plt.figure(figsize=(8, 6))
disp.plot(cmap=plt.cm.Blues, values_format='d')
plt.title('Confusion Matrix')
plt.show()

# Calculate Mean Squared Error (MSE)
mse = mean_squared_error(y_test, y_pred_proba)
print(f"Mean Squared Error (MSE): {mse:.4f}")

# Calculate Root Mean Squared Error (RMSE)
rmse = np.sqrt(mse)
print(f"Root Mean Squared Error (RMSE): {rmse:.4f}")

# Sort values for a smooth line plot
X_test_sorted, y_pred_proba_sorted = zip(*sorted(zip(X_test['Age'], y_pred_proba)))

# Plotting Age vs Outcome using Logistic regression probabilities
plt.figure(figsize=(10, 6))
plt.plot(X_test_sorted, y_pred_proba_sorted, color='red', label='Logistic Regression')
plt.scatter(X_test['Age'], y_test, color='blue', alpha=0.5, label='Actual Data Points')
plt.xlabel('Age')
plt.ylabel('Outcome Probability')
plt.title('Logistic Regression with SGD: Age vs Outcome')
plt.legend()
plt.show()
```


	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1


Accuracy: 0.3594

Precision: 0.3594

Recall: 1.0000

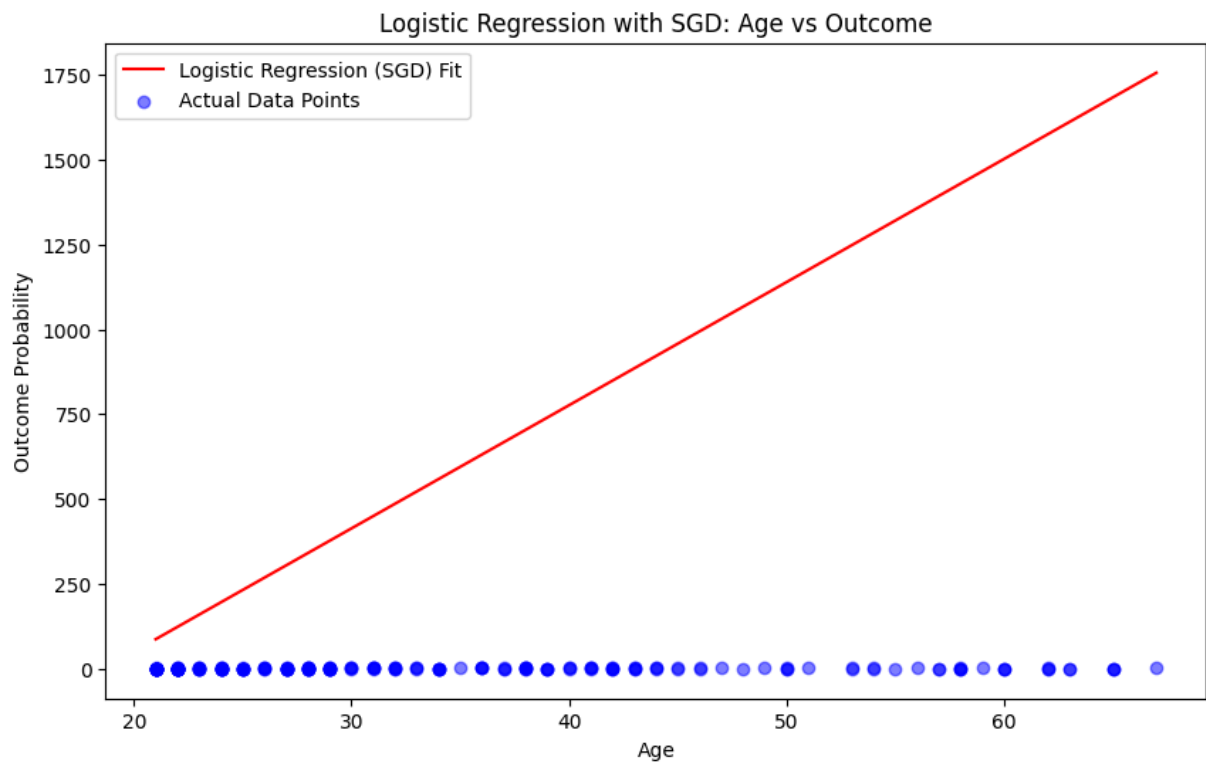
F1 Score: 0.5287

<Figure size 800x600 with 0 Axes>



Mean Squared Error (MSE): 523140.6782

Root Mean Squared Error (RMSE): 723.2846



Naive Bayes Classifier

In []: